

# CS061 - Lab 01

## Welcome to LC3

### 1 High Level Description

Today's lab will cover the basics of how to use the LC3 Assembler/Simulator, and write a basic program in LC3.

The text for today's lab is rather long.

Don't panic! Most labs will not be this verbose! There is much introductory material to cover.

### 2 Objectives for This Week

1. Learn the basic skeleton for any LC3 program file
2. Introduction to the basics of LC3 programming
3. Exercise 0: Hello World!
4. Exercise 01: Implement, run, and inspect a simple program yourself
5. Uh, ok... now what? (aka: What to study for next time)

### 3 Alright, let's rock these objectives!

#### 3.1 Basic skeleton for any LC3 program

Similar to C++, every LC3 program has a certain amount of overhead necessary for making everything work properly. In C++, you may have something like the following:

```
//  
// File: main.cpp  
// Description: A simple hello world-esque program  
// Author: Sean Foley  
// Date created: 12/19/2009  
// Date last modified: 12/19/2009  
//  
#include <iostream>  
using namespace std;  
  
int main()  
{  
    cout << "Egads, this is the most fun I've had in years!" << endl;  
}
```

This source code is passed to a compiler, which ultimately produces a machine language executable.

In LC3, a basic program will have the following look:

```
;  
; Foley, Sean  
; Login: sfoley (sfoley@cs.ucr.edu)  
; Section: xxx  
; TA: Sean Foley  
; Lab 01  
;  
.orig x3000      ; **ALL** your programs will start at x3000  
; Instructions  
    ; LC3 instruction code goes here  
; Local Data  
    ; pseudo-ops for hard-coding data go here  
.end    ; .end is like the "}" after main() in C++. It means "no more code to  
compile!"
```

If you haven't figured it out by now, in LC3, comments are denoted with semi-colons. They can be on a line by themselves or on the same line as actual code (at the end of the line)

This source code is passed to an assembler - a much simpler beast than a compiler, but just like a compiler it produces a machine language executable.

**NOTE: If you don't like losing points, make sure you include the above header in ALL your LC3 lab assignments, programming assignments, and lab practicals.**

## 3.2 Introduction to the basics of LC3 programming

An LC3 program consists of three basic elements: **labels**, **pseudo-ops**, and **instructions**.

**Labels** are simply “aliases” for memory addresses. These are really, really useful because they relieve the programmer of the responsibility of figuring out which line of code is at which memory address, and also allow us to use symbolic names to access memory locations containing data.

**Pseudo-ops** tell the assembler how to set things up before starting to translate the source code into machine language. They are just like compiler directives in C++, like `#include`.

For example, the pseudo-op **.FILL** writes a hard-coded value into memory (i.e. RAM, or system memory). For example, the two **.FILL** pseudo-ops in Table 1 would write the value 6 into memory at x3005 and the value 12 into memory at x3006 before the program begins execution. Note that you don't need to know the memory addresses corresponding to the labels – that's the assembler's job.

Table 1: Labels and Pseudo-ops

<u>Memory address</u>	<u>Label</u>	<u>Pseudo-Op</u>	<u>Hard-coded value</u>
x3005	DEC6	.FILL	#6
x3006	DEC12	.FILL	#12

*(The '#' before the numbers means decimal – i.e base 10; 'x' means hexadecimal – i.e. base 16).*

Now we can refer to these memory locations (and therefore access the data they contain) using the names DEC6 and DEC12.

We will learn about another pseudo-op **.STRINGZ** in a moment.

**Instructions:** There are 15 LC3 instructions, in three distinct categories:

- Arithmetic / Logic operations (adding, boolean operations, etc)
- Data movement (moving data between RAM and local registers)
- Program control (changing the order of execution of code - if statements, for loops, etc)

**Following are one or two instructions from each of these categories.**

### 3.2.1 Arithmetic / Logic Operations

These instructions are used for manipulating values and doing calculations. For example: Adding two numbers together, AND'ing two numbers together, or taking the logical NOT of a number.

The LC-3 has 8 general purpose, 16-bit **local registers** (analogous to variables in C++), named R0 through R7, that can be used for performing such operations.

The **ADD** instruction (integer data type)

This instruction adds the contents of two registers, and writes the result to a third. It comes in two flavors:

See the [ADD tutorial](#) (on the main course page under “LC3 instruction tutorials”) for example code, descriptive visuals, and theoretical content.

### 3.2.2 Data movement

Data movement instructions are used when you want to copy a value from a memory location into a register (called “loading”), or copy a value from a register to a memory location (“storing”).

You will be doing this a **lot**, since the LC3 can operate only on values in registers, not in memory.

The **LEA** instruction (“Load Effective Address”)

This instruction translates a label – the name you give to a memory location – back into the memory address it stands for, and stores that address into the specified register.

There is no actual data movement here – just the decoding of a label. We will see later how it is used to support actual data movement instructions.

See the [LEA tutorial](#) (on the main course page under “LC3 instruction tutorials”) for example code, descriptive visuals, and theoretical content.

The **LD** instruction (“Load Direct” - probably the most frequently used instruction!)

This instruction copies the contents of a specified memory location into a register.

The location is specified by its label.

See the [LD tutorial](#) (on the main course page under “LC3 instruction tutorials”) for example code, descriptive visuals, and theoretical content.

### 3.2.3 Program control

These instructions are used in control structures such as **BRANCHING**, **LOOPING**, and **FUNCTION CALLS** (aka: **Subroutine Calls**).

The **BR** instruction (“Conditional Branch”)

Ordinarily, a program starts at the first instruction and executes one line of code after another until it reaches the end (“sequential execution”).

The **Branch** instruction can be used to alter this flow of execution based on a *condition*.

**All** the control structures you are familiar with from C++ (if, for, while, do, etc) can be built using this one instruction.

The conditional branch instruction has three modifiers: {**n**, **z**, **p**}.

It examines the value of the last modified register (“**LMR**”) and jumps to the memory location specified by the label **IF AND ONLY IF** any one of the specified conditions is met.

See the [BR tutorial](#) (on the main course page under “LC3 instruction tutorials”) for example code, descriptive visuals, and theoretical content.

## 3.3 **Exercise 0: HELLO WORLD** (*aaargh!! not “hello world” again?!?!)*

As always, the very first program we will write will simply output the message “Hello World!”

Though simple, it will illustrate several of the points above, plus some things you will fully understand only later on (*exactly like when you first encountered `cout` in C++*):

- We will store the string as an array of ASCII characters in memory, starting at an address labelled “message”, using the pseudo-op **.STRINGZ**  
This pseudo-op stores a string in memory, one character per memory address, with a zero (“the null character” or \0) after the last character.  
You will recognize this as a “c-string”, i.e. a null-terminated character array.  
`shrug .STRINGZ "whatever"`  
will write the characters 'w' 'h' 'a' 't' 'e' 'v' 'e' 'r' followed by x0000, to the nine memory locations starting at the label shrug.
- We will use the LEA instruction to translate the label into an actual address, which will be stored in R0.
- We will invoke an i/o subroutine called PUTS, which uses R0 to locate the start of the string to be output (it knows when to stop because the last value in the char array is \0).

[illegible]

Now write this program for yourself:

1. Create a new file called **helloworld.asm** in whatever editor you like (vim, gedit, emacs)
2. Type in the code above
3. Save and close the file
4. Open your program in the LC3 assembler & simulator by typing **simpl helloworld.asm** on the command line.  
This will open two windows: one showing the assembled code, the other for input/output  
Make sure you can see both.
5. Hit the **Run** button at the bottom of the simulator window, and observe the worlds “Hello World!” magically appear in the console window!

### 3.4 Exercise 01: A “real” program

Now Hello World is out of the way, we can write a program that actually does some processing: it will multiply a number, which we will store in advance in memory, by six.

We will also take this opportunity to introduce the style you **MUST** use for all your LC-3 programs.

```
;-----  
; Foley, Sean  
; Login: sfoley (sfoley@cs.ucr.edu)  
; Section: xxx  
; TA: Sean Foley  
; Lab 01  
;-----  
.orig x3000  
;-----  
; Instructions  
;-----  
LD R1, DEC_0          ; R1 <-- #0  
LD R2, DEC_12         ; R2 <-- #12  
LD R3, DEC_6          ; R3, <-- #6  
  
DO_WHILE_LOOP  
    ADD R1, R1, R2      ; R1 <-- R1 + R2  
    ADD R3, R3, #-1     ; R3 <-- R3 - #1  
    BRp DO_WHILE_LOOP  ; if (R3 > 0): goto DO_WHILE_LOOP  
END_DO_WHILE_LOOP  
  
HALT                  ; halt program (like exit() in C++)  
;-----  
; Local data  
;-----  
DEC_0    .FILL    #0    ; put #0 into memory here  
DEC_12   .FILL    #12   ; put #12 into memory here  
DEC_6    .FILL    #6    ; put #6 into memory here  
  
.end
```

First comes the header, which you ***MUST*** place at the start of every program.

Next, we have the line ".orig x3000"

This indicates where in memory the code will reside. The LC3 assembler requires all programs to be loaded at or above x3000: all your programs ***must*** load the "main" routine to x3000  
*(just try something else and see what happens!! On second thoughts – don't!)*

The next three lines load hard-coded values from memory into registers R1, R2, and R3 respectively. R1 is an accumulator, so we initialize it to 0; R2 holds the number to be multiplied; and R3 is used as a loop counter (to keep track of how many times the loop will execute – i.e. a "counter-controlled" loop).

Note how the data is stored in a block separated from the code, ***after*** the HALT instruction – all your programs must do the same.

*Can you see how this program works?*

Now that you have had a nice lengthy overview of some basic LC3 instructions and seen how to use the simpl assembler/simulator, it is time for you to implement the program yourself.

Proceed as in exercise 0; call your file **lab1.asm**

This time, instead of hitting Run, step through each line of code using the **Step** button in the simulator until you reach the HALT instruction. Notice how the value of R1 keeps going up by 12 until it reaches *(guess what??) 72*  
*(Now, how did it know to stop at just the right number??)*

Finally, stare in (pick one) {wonder, horror, bewilderment} at the screen

### 3.5 Uh, ok... now what? (aka: What to Study for Next Time)

For next lab, make sure you understand:

- Basic required headers for all LC3 programs in this course
- That all programs start at x3000
- What labels are and how to use them
- How to use the .FILL and .STRINGZ pseudo-ops
- The LD, LEA, BR, HALT instructions
- The PUTS output routine (aka: Trap x22)
- How to create, save, run, and step through an LC3 program
- The logic of the program in Exercise 01