

Introduction to Computing Systems

from bits & gates to C & beyond

Chapter 5

The LC-3 Instruction Set Architecture

- *ISA Overview*
- *Operate instructions*
- *Data Movement instructions*
- *Control Instructions*
- *LC-3 data path*

LC-3 ISA Overview

• Memory organization

- Address space: $2^{16} = 64k$ locations
- Addressability: Word
- Word size = 2 bytes
 - \Rightarrow *total memory* = $64k \times 2 = 128$ kbytes

• Registers

- 8 x 16 bit General Purpose Registers: R0 - R7
- 3 x 1 bit Condition Code Registers: N, Z, P
- Some special purpose registers (more later)

• Instructions

- 16 bit instructions, with 4 bit opcodes
- Native Data Type: only 2's complement integer
- Addressing Modes: Immediate, Register, Direct, Indirect & Base+Offset

LC-3 Instructions

- **Operate**

- **Manipulate data directly - arithmetic or bitwise logic**
 - *ADD, AND, NOT*

- **Data Movement**

- **Move data between memory and registers**
 - *LD, LDI, LDR, LEA, ST, STI, STR*

- **Control**

- **Change the sequence of instruction execution**
 - *BR, JMP/RET, JSR/JSRR, TRAP, RTI*

Instruction Construction

•Two main parts

- Opcode: specifies what the instruction does.
- Operand(s): what the instruction acts on.
- Instruction sets can be complex or simple (CISC, RISC), single-word or multi-word.

•LC-3

- Single word (16 bit) instructions.
- 4-bit opcode => 16 instructions (very simple set!)
- remaining 12 bits specify operand(s), according to the addressing mode proper to each instruction.

LC 3 Instructions

• LC-3 Instruction word: 16 bits

• Opcode

- *IR[15:12]: 4 bits allow 16 instructions*
- *specifies the instruction to be executed*

• Operands

- *IR[11:0] contains specifications for:*
 - Immediate value: 5 bits
 - Registers: 8 GPRs (i.e. require 3 bits for addressing)
 - Address Generation bits: “Offset” (11 or 9 or 6 bits - more later)

• Examples

ADD DR, SR1, SR2 ; DR \leftarrow (SR1) + (SR2)
[15:12] [11:9] [8:6] [2:0]

LDR DR, BaseR, Offset ; DR \leftarrow Mem[(BaseR) + Offset6]
[15:12] [11:9] [8:6] [5:0]

Addressing Modes

Note: the effective address (ea) is the memory location of the operand

- The LC-3 supports five addressing modes:
 - the operand is located:
 - *in the instruction itself* (*immediate*)
 - *in a register*
 - *in memory:*
 - the ea is encoded in the instruction (*direct, or PC-relative*)
 - a pointer to the ea is encoded in the instruction (*indirect*)
 - a pointer to the ea is stored in a register (*relative, or base+offset*)

Operate Instructions

- **NOT (unary operator)**

- destination register in IR[11:9] and a single source register in IR[8:6].
- bits IR[5:0] are all 1s.

- **ADD & AND (binary operators)**

- destination register in IR[11:9], one source register in IR[8:6]
- other source:
 - *immediate addressing mode:*
 - if bit IR[5] = 1, bits IR[4:0] specify the other source number directly, as a 5 bit 2's complement integer, which is sign extended (SEXT) to 16 bits.
 - *register addressing mode:*
 - if bit IR[5] = 0, bits IR[2:0] specify a register for the second source
 - bits IR[4:3] = 0

Immediate & Register Operands

• Immediate

opcode		operands			
[15:12]	[11:9]	[8:6]	[5]	[4:0]	
ADD	DR	SR1	1	imm	

- If bit 5 = 1, the value in IR[4:0] (“immediate”) is *sign extended (SEXT)* to 16 bits and added to the contents of the source register SR1 (IR[8:6]).

• Register

opcode		operands			
[15:12]	[11:9]	[8:6]	[5]	[2:0]	
ADD	DR	SR1	0	SR2	

- if bit 5 = 0, the contents of source register SR2 (IR[2:0]) are added to the contents of source register SR1 (IR[8:6]).
- In both cases, the result goes to the destination register DR (IR[11:9]).

NOT: Bitwise Logical NOT

- **Assembler Inst.**

NOT DR, SR ; DR <= NOT (SR)

- **Encoding**

1001 DR SR 111111

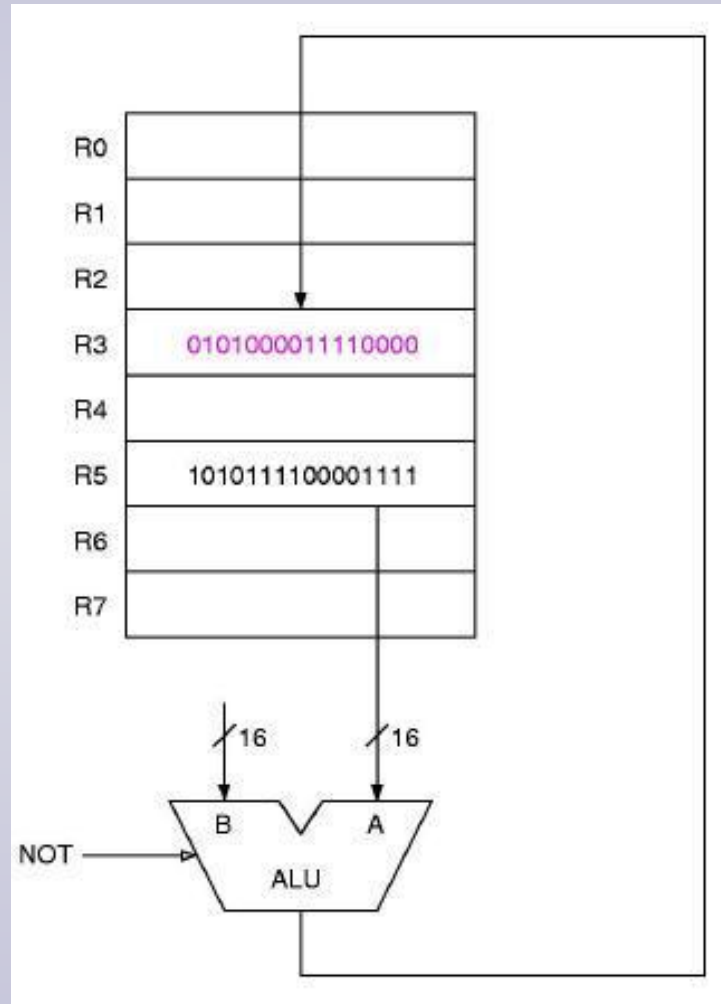
- **Example**

NOT R2, R6

- **Note: Condition codes are set.**

NOT data path

NOT R3, R5



ADD: Two's complement 16-bit Addition

- **Assembler Instruction**

ADD DR, SR1, SR2 ; DR <= (SR1) + (SR2) (*register addressing*)

ADD DR, SR1, imm5 ; DR <= (SR1) + Sext(imm5) (*immediate addressing*)

- **Encoding**

0001 DR SR1 0 00 SR2

0001 DR SR1 1 imm5

- **Examples**

ADD R1, R4, R5

ADD R1, R4, #-2

- **Note: Condition codes are set**

AND: Bitwise Logical AND

- **Assembler Instruction**

AND DR, SR1, SR2 ; DR <= (SR1) AND (SR2)

AND DR, SR1, imm5 ; DR <= (SR1) AND Sext(imm5)

- **Encoding**

0101 DR SR1 0 00 SR2

0101 DR SR1 1 imm5

- **Examples**

AND R2, R3, R6

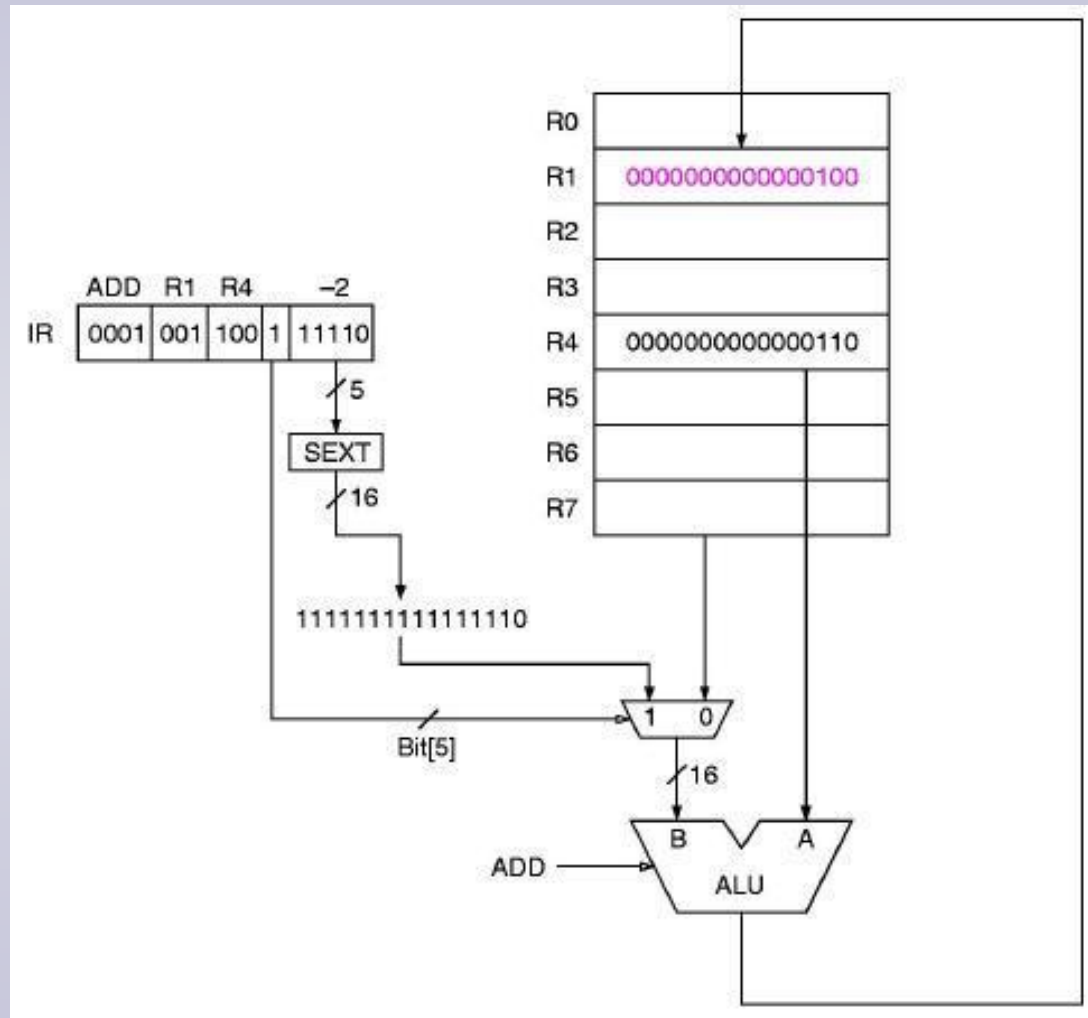
AND R2, R2, #0 ; Clear R2 to 0

Question: *if the immediate value is only 6 bits, how can it mask the whole of R2?*

- **Note:** Condition codes are set.

ADD data path

ADD R1, R4, # -2



Data Movement Instructions - 1

- **Move Data**

- from register to memory => store
 - *nominated register is Source*
- from memory to register => load
 - *nominated register is Destination*
- The LC-3 cannot move data from memory to memory
- also to/from I/O devices (later)

- **LC-3 Load/Store Instructions**

- LD, LDI, LDR, LEA, ST, STI, STR
- Format:

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
opcode				DR or SR			Address generator bits								

Note: The field 'PCoffset9' in the following slides is IR[8:0]

Data Movement Instructions - 2

- **LC-3 Load/Store Addressing modes:**

- ***immediate: LEA***

No Effective Address (EA) calculation; the Sign Extended Addr. Generator is added to the current value of the Program Counter - i.e.

$$DR \leq (PC) + \text{SEXT}(IR[8:0])$$

- ***direct or PC-Relative: LD & ST***

The EA is the Sign Extended Addr. Generator added to the current value of the Program Counter - i.e.

$$EA = (PC) + \text{SEXT}(IR[8:0])$$

$$DR \leq \text{Mem}[(PC) + \text{SEXT}(IR[8:0])]$$

- ***indirect: LDI & SDI***

$$EA = \text{Mem}[(PC) + \text{SEXT}(IR[8:0])]$$

$$DR \leq \text{Mem}[\text{Mem}[(PC) + \text{SEXT}(IR[8:0])]]$$

- ***base+offset: LDR & STR (BaseReg is specified by IR[8:6])***

$$EA = \text{BaseReg} + \text{SEXT}(IR[5:0])$$

$$DR \leq \text{Mem}[(\text{BaseReg}) + \text{SEXT}(IR[5:0])]$$

Memory Addressing Modes

- Direct addressing (PC-Relative)



- effective address = (PC) + SEXT(IR[8:0])
- operand location must be within approx. 256 locations

of the instruction

- *actually between +256 and -255 locations of the instruction being executed (why?)*

Memory Addressing Modes - 2

- Indirect addressing



- Same initial mechanism as direct mode (i.e. PC-Relative), but the calculated memory location now contains the *address of the operand*, (i.e. the ea is indirect):


pointer address = (PC) + SEXT(IR[8:0])

effective address = Mem[(PC) + SEXT(IR[8:0])]

- *Note that the memory has to be accessed twice to get the actual operand.*

Memory Addressing Modes - 3

- Base+Offset addressing



- **effective address = (BaseRegister) + offset**
 - *sign extend (SEXT) the 6 bit offset ([5:0]) to 16 bits*
 - *add it to the contents of the Base Register ([8:6])*
- **differences other addressing modes:**
 - *base+offset field is 6 bits, PC-Relative offset field is 9 bits.*
 - *base+offset can access any location in memory, PC-Relative offset only within +/- 256 locations of the instruction.*
 - *base+offset accesses memory once, Indirect mode accesses memory twice.*

LD: Load Direct

- Assembler Inst.

LD DR, LABEL ; DR <= Mem[LABEL]

- Encoding

0010 DR PCOffset9 ; EA = (PC) + SEXT(PCOffset9)
= (PC) + SEXT(IR[8:0])

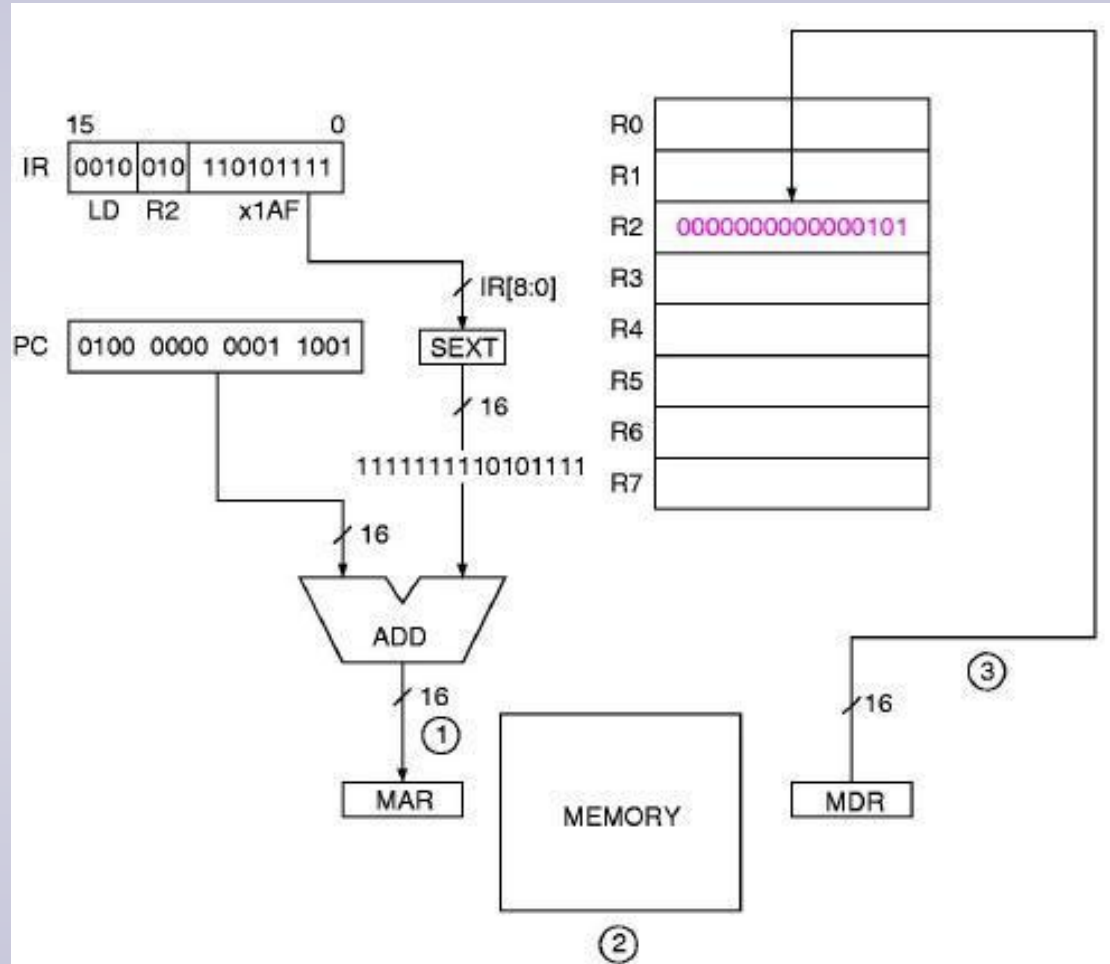
- Examples

LD R2, param ; R2 <= Mem[param]

**Notes: The LABEL must be within +256/-255 lines of the instruction.
Condition codes are set.**

LD data path

LD R2, x1AF



LDI: Load Indirect

- Assembler Inst.

LDI DR, LABEL ; DR <= Mem[Mem[LABEL]]

- Encoding

1010 DR PCoffset9 ; EA = Mem[(PC) + SEXT(PCoffset9)]
= Mem[(PC) + SEXT(IR[8:0])]

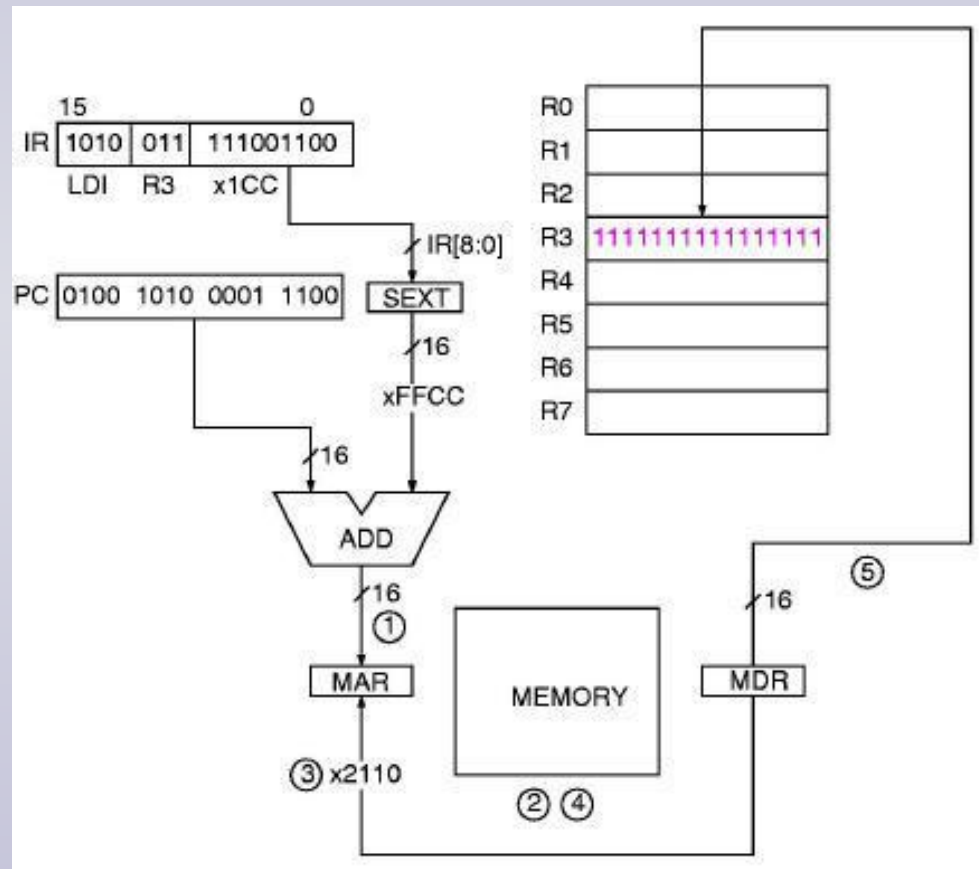
- Examples

LDI R2, POINTER ; R2 <= Mem[Mem[POINTER]]

**Notes: The LABEL must be within +256/-255 lines of the instruction.
Condition codes are set.**

LDI data path

LDI R3, x1CC



LDR: Load Base+Offset

- Assembler Inst.

LDR DR, BaseR, offset ; DR <= Mem[(BaseR) + SEXT(IR[5:0])]

- Encoding

0110 DR BaseR offset6

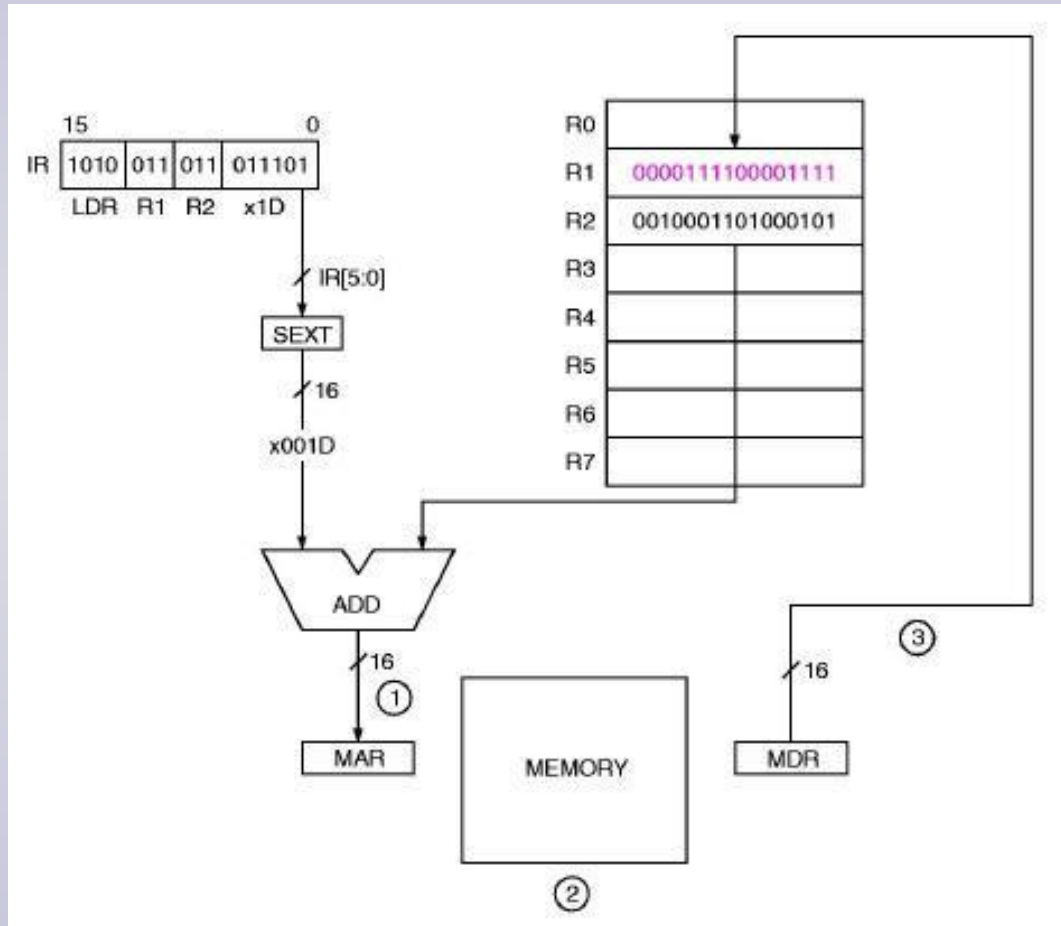
- Examples

LD R2, R3, #15 ; R2 <= Mem[(R3)+15]

**Notes: The 6 bit offset is a 2's complement number, so range is -32 to +31.
Condition codes are set.**

LDR data path

LDR R1, R2, x1D



LEA: Load Effective Address

- Assembler Inst.

LEA DR, LABEL ; DR <= LABEL

- Encoding

1110 DR offset9 i.e. address of LABEL = (PC) + SEXT(offset9)

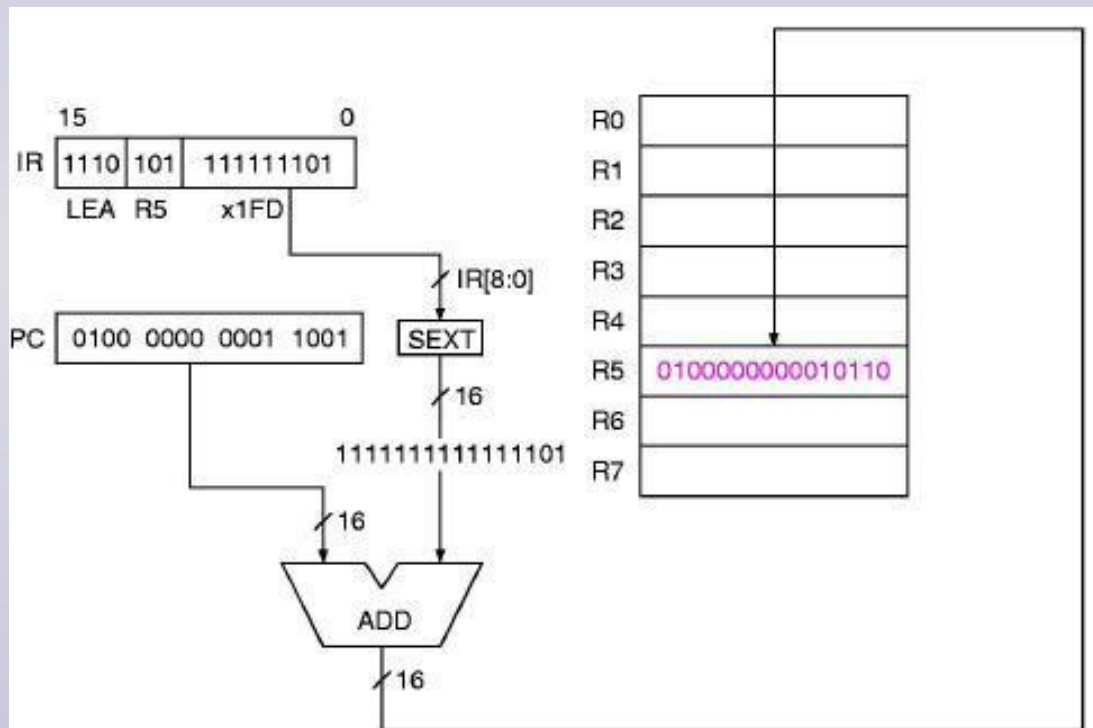
- Examples

LEA R2, DATA ; R2 gets the address of DATA = (PC) + Sext(IR[8:0])

**Notes: The LABEL must be within +/- 256 lines of the instruction.
Condition codes are set.**

LEA data path

LEA R5, # -3



ST: Store Direct

- Assembler Inst.

ST SR, LABEL ; [LABEL] <= (SR)
; The value stored in SR is copied to the
; memory address identified as LABEL

- Encoding

0011 SR offset9

- Examples

ST R2, VALUE ; [VALUE] <= (R2)

**Notes: The LABEL must within +/- 256 lines of the instruction.
Condition codes are NOT set.**

STI: Store Indirect

- Assembler Inst.

STI SR, LABEL ; [Mem[LABEL]] <= (SR)
; The value stored in SR is copied to the
; memory address pointed to by LABEL

- Encoding

0011 SR offset9

- Examples

STI R2, POINTER ; [Mem[POINTER]] <= (R2)

**Notes: The LABEL must be within +/- 256 lines of the instruction.
Condition codes are NOT set.**

STR: Store Base+Offset

- Assembler Inst.

STR SR, BaseR, offset6 ; [(BaseR)+SEXT(offset6)] <= (SR)
; EA = (BaseR) + SEXT(offset6)
; The contents of SR are copied to EA

- Encoding

0111 SR BaseR offset6

- Examples

STR R2, R4, #15 ; [(R4) +15] <= (R2)

Notes: The offset is sign-extended to 16 bits.
Condition codes are not set.

Addressing Examples

What is the EA for the following instructions?

Given:

(PC) = x2081, (R6) = x2035, LOC = x2044, Mem[x2044] = x3456

LDI R2, LOC

Encoding:

1010 010 1 1100 0011

Indirect addressing:

EA = Mem[?] = ?

LDR R1, R6, #12

Encoding:

0110 001 110 00 1100

Base+Offset addressing:

EA = ? = ?

ADD R1, R3, R2

Register addressing:

DR = R1, SR1 = R3, SR2 = R2

DR <= ?

ADD R5, R1, #15

Immediate addressing:

DR = R5, SR1 = R1, S2 = 15

DR <= ?

LD R1, LOC

Direct addressing:

DR <= ?

Control Instructions

- **Change the Program Counter**
 - Conditionally or unconditionally
 - Store the original PC (subroutine calls) or not (go-to)
- **LC-3 Control Instructions**
 - **BRx, JMP/RET, JSR/JSRR, TRAP, RTI**
 - *BRx uses PC-Relative addressing with 9-bit offset*
 - *JSR uses PC-Relative addressing with 11-bit offset*
 - *JMP/RET & JSRR use base+offset addressing with zero offset*
 - *we'll deal with TRAP & RTI later*

BR: Conditional Branch

- Assembler Inst.

BRx LABEL ; PC <= LABEL iff condition evaluates to true
where x = n, z, p, nz, np, zp, or nzp
Branch to LABEL iff the selected condition code are set

- Encoding

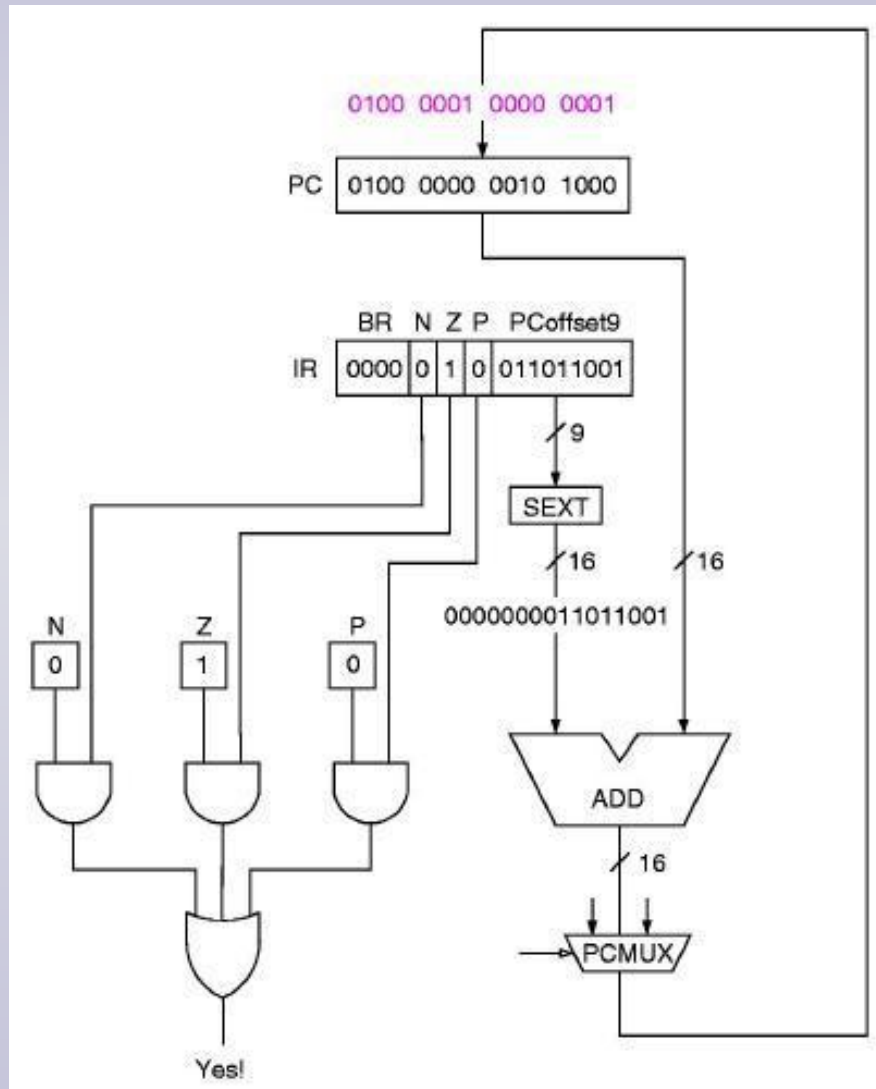
0000 n z p PCoffset9 ; PC <= (PC) + Sext(IR[8:0]) if condition

- Examples

BRzp LOOP ; branch to LOOP if previous op returned zero or positive.

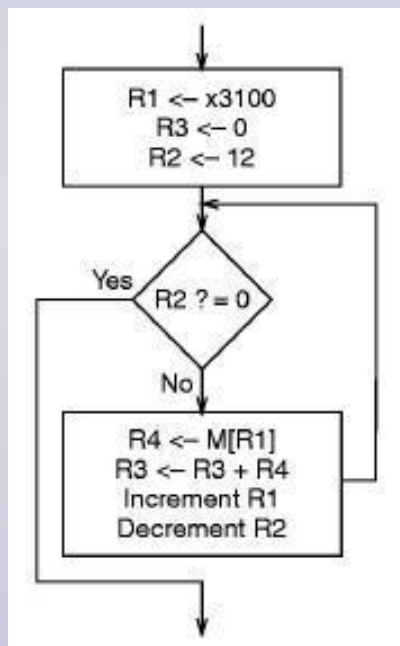
BR data path

BRz x0D9

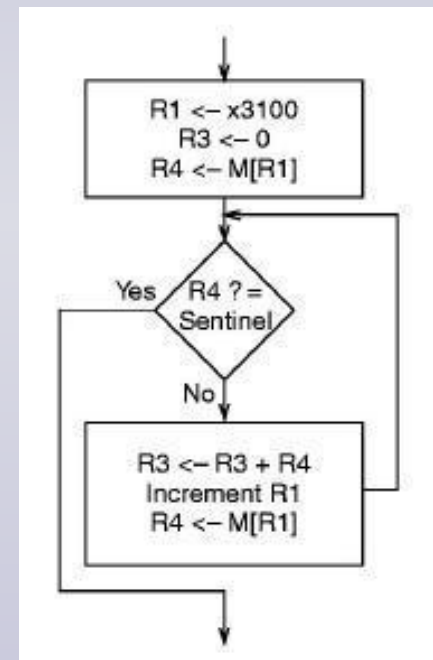


Building loops using BR

Counter control



Sentinel control



JMP: Jump or Go To

- Assembler Inst.

JMP BaseR ; PC <= (BaseR)

Take the next instruction from the address stored in BaseR
(this is the same memory addressing mode as LDR & STR, but with no offset)

- Encoding

1100 000 BaseR 00 0000

- Example

JMP R5 ; if (R5) = x3500, the address x3500 is written to the PC

TRAP Instruction

- Used to invoke an operating system service.
- Trap vector table: a list of locations of the service call routines.
- TRAP has one operand, the trap vector:
PC is set to the value stored at that location of the vector table.
- Some special trap vectors:
 - * x20: input a character from the keyboard
 - * x23: input a character from the keyboard, with prompt & echo
 - * x21: output a character to the console display
 - * x25: halt the program
- More details later

TRAP: Invoke a system routine

- Assembler Inst.

TRAP trapvec ; step 1: R7 <= (PC)

; step 2: PC <= Mem[Zext(IR[7:0])]

Note we ZERO extend the trapvec, not SIGN extend

- Encoding

1111 0000 trapvect8

- Example

TRAP x22 ; invoke the BIOS routine PUTS

Meaning: R7 <= (PC) (*for eventual return from subroutine*)

PC <= Mem[x0022]

i.e. PC gets the address stored at entry trapvect8 in the trap vector table.

Data Path - 1

•Global Bus

- 16-bit, data & address
- connects all components
- is shared by all

•Memory

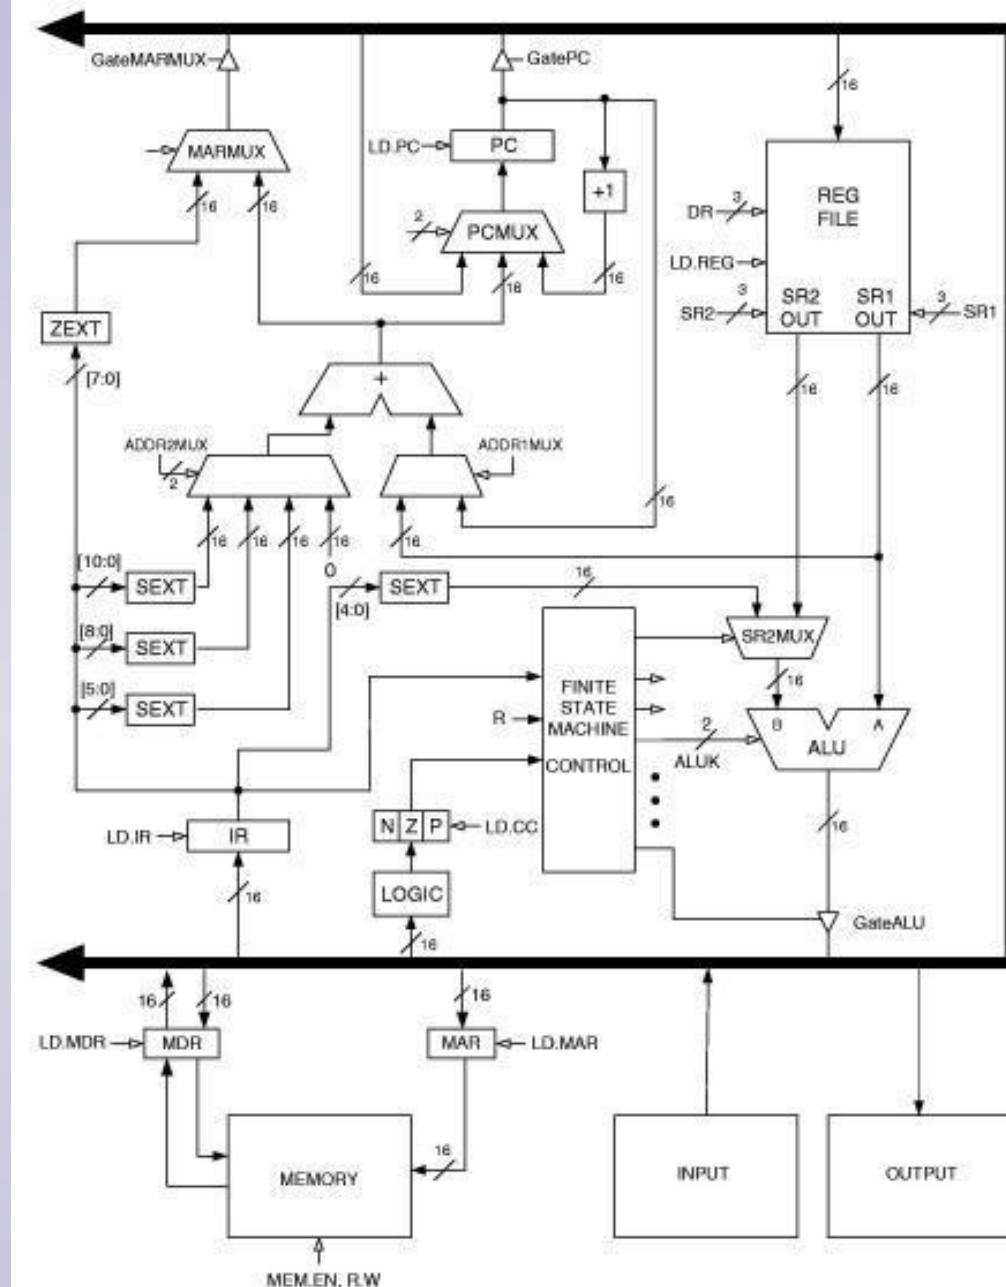
▪ Memory Address

Register: MAR

- address of location to be accessed

▪ Memory Data Register: MDR

- data loaded or to be stored



- **ALU & Registers**

- PC & PCMUX

-
- The diagram illustrates the internal architecture of the MARS computer. Key components and their interconnections include:
- Control and State:** A **FINITE STATE MACHINE** and **CONTROL** unit manage the system. A **LOGIC** block provides status signals **N**, **Z**, and **P** to the **PC** and **PCMUX**. The **PC** (Program Counter) is updated via **LD.PC** and **GatePC**. The **PCMUX** selects between the **PC** and a constant **+1** to update the **PC**.
 - Instruction and Register File:** The **IR** (Instruction Register) is loaded via **LD.IR**. The **REG FILE** (Register File) has **SR1** and **SR2** ports for reading and writing, controlled by **LD.REG** and **SR2** signals. It also has **DR** (Data Register) and **SR1 OUT** ports.
 - Arithmetic and Logic:** The **ALU** (Arithmetic Logic Unit) performs operations on data from the **REG FILE** and **PCMUX**. It is controlled by **GateALU** and **ALUK** signals. The **ALU** output is connected to the **OUTPUT** block.
 - Memory and I/O:** The **MEMORY** block is connected to the **MDR** (Memory Data Register) via **LD.MDR** and **MEM.EN, R.W** signals. The **MAR** (Memory Address Register) is loaded via **LD.MAR** and provides the address for the **MEMORY**. The **INPUT** block provides data to the **MEMORY**.
 - Multiplexers and Registers:** The **MARMUX** selects between the **PC** and the **ALU** output to update the **PC**. The **PCMUX** selects between the **PC** and the **+1** constant to update the **PC**. The **SR2MUX** selects between the **REG FILE** and the **ALU** output to update the **REG FILE**.

- **MARMUX**

-